

Replex: A Scalable, Highly Available Multi-Index Data Store

By Amy Tai et al. - *Princeton University*

Presented by Camilla Stormoen

Overview

- Motivation
- System Design
- Hybrid Replex
- Evaluation
- Conclusion

Motivation

- Need for scalable, high-performance data stores ➡ NoSQL databases


How do we put back the secondary indices, without compromising scalability, availability and performance?

- Replex
 - enables efficient querying on multiple keys
- Hybrid Replexes
 - enable rich design space for trading off steady-state performance with faster recovery

System Design

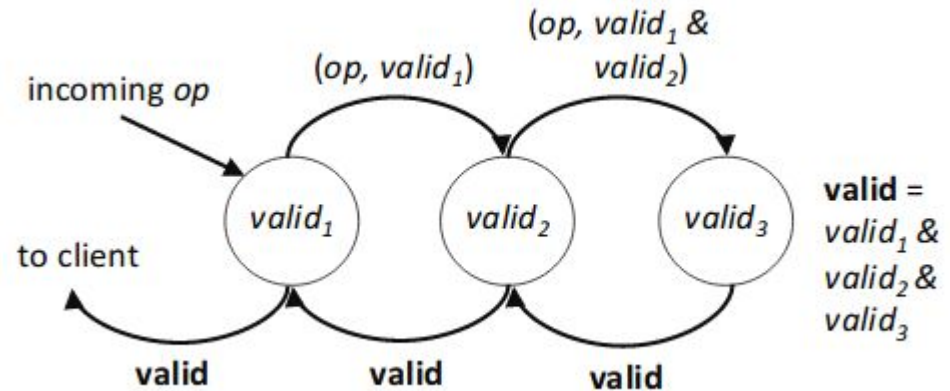
- Data Model and API
 - Stores data in the form of RDBMS tables
 - Table has schema with columns
 - Table-column primary key ➡ index of table

- Data Partitioning with Replexes
 - Fast query: Partitioned by index
 - Replex stores a table and shards the rows across multiple partitions
 - Chain replication



System Design: Replication Protocol Failure

- Individual replexes can have requirements
 - Can both be valid and invalid depending on the replex
- Chain replication includes a consensus protocol




System Design: Amplification

- Two concerns:
 1. How to reconstruct the failed partition
 2. How to respond to queries that would have been serviced by the failed partition.
- Partition fails: read must broadcast to every partition and iterate through the entire local storage of partition
- Within a threshold: introduce f replicas with same sharding function
 - - Cost is storage and network overhead

Hybrid Replexes

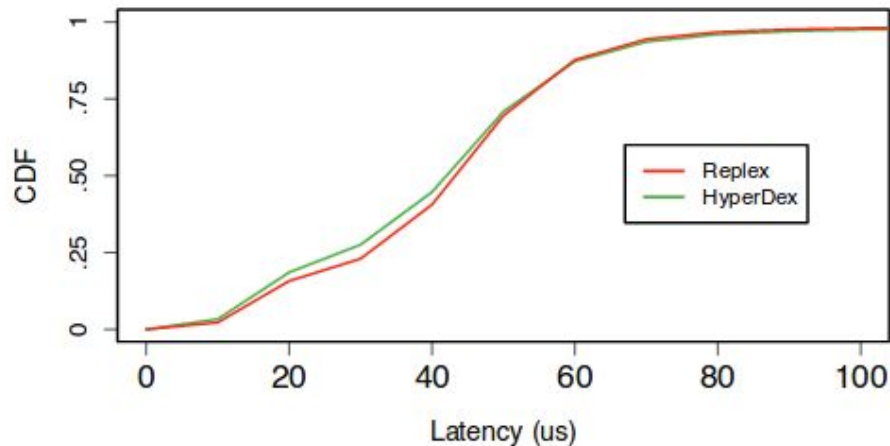
- Increase failure resilience of any number of replexes
- 2-Sharing
 - Two partitions share data if there is a path between them
 - P shares data with two partitions in hybrid replex and four in replex A
 - Read redirect to hybrid replex is faster compared to redirect to hybrid A




Evaluation: Steady-State Performance

- *How does Replex's design affect steady-state index performance?*
- Single hybrid replex (*Replex-3* - 3 replexes, tolerant to 2 failures)

- Read latency



- Insert latency



Evaluation: Failure evaluation

- *How do hybrid replexes enable superior recovery performance?*
- 12 virtual partitions per subspace/replex.
- Table with primary and secondary index. Split 50:50 reads
- Recovery time: reduce by 2-3x
- Recovery throughput: Replex-2 minimal throughput




Table 1: 10 millions rows of size 100 bytes records

System	Recovery Time (s)	Recovery Throughput (op/s)
Replex-2	50 ± 1	$18,989 \pm 1,883$
Replex-3	60 ± 1	$65,780 \pm 3,839$
HyperDex	105 ± 17	$34,697 \pm 19,003$

Table 2: 1 million rows of size 1 KB records

System	Recovery Time (s)	Recovery Throughput (op/s)
Replex-2	6.7 ± 0.57	$70,084 \pm 5,980$
Replex-3	8.7 ± 0.56	$110,280 \pm 11,232$
HyperDex	20.0 ± 2.65	$127,232 \pm 85,932$

Conclusion

- Able to query data by more than just a single key
- Replex's steady-state performance 76% better than HyperDex
- Multi-Index, scalable, highly-available NoSQL data store is possible and a better choice

Questions?